



The Complete Guide to Measuring What Matters in Generative RAG AI Systems

9 min read · Oct 14, 2025



Dewasheesh Rana

Follow

Listen

Share

More

Introduction — From Smart to Measurable AI

Retrieval-Augmented Generation (RAG) is at the heart of today's GenAI systems — from copilots that debug code to assistants that explain financial transactions or resolve incidents.

But here's the truth:

A RAG pipeline that *sounds smart* isn't necessarily *doing its job right*.

To build AI systems that are reliable, explainable, and production-grade, you must **measure** their behavior — not just their accuracy.

In this article, we'll dive into the full landscape of **RAG and Agentic RAG metrics** — starting from the basics like Precision and Recall, then moving to advanced measures like nDCG, Faithfulness, Context Precision, and RAGAS — with detailed numerical examples anyone can understand.

By the end, you'll know exactly how to evaluate a GenAI system the same way you'd test a financial algorithm or cloud microservice: with hard, measurable evidence.

Understanding the RAG Pipeline

Before we discuss metrics, let's recall the architecture:

User Query → Retriever → Context → Generator → Final Answer

The **Retriever** fetches relevant data (from vector DBs like Qdrant, OpenSearch, Pinecone, etc.).

The **Generator** (LLM) reads this context and produces the answer.

Together, they form the **RAG pipeline**.

Agentic RAG extends this with **autonomous decision-making** — where agents (like CrewAI) can:

- invoke tools,

[Open in app](#) ↗

≡ **Medium**



- execute corrective actions.

So, metrics matter even more — because now, your AI is *acting* as well as *thinking*.

PART 1 — Retrieval Metrics

How good is your “librarian”?

Retrieval metrics evaluate the quality of the context your LLM is fed.

If you feed garbage, you’ll get confident-sounding garbage back.

Example

Query: “*Why did the Airflow DAG fail yesterday?*”

Retriever fetches 5 log snippets:

Rank	Log Snippet	Relevant?
1	Timeout error on operator	✓
2	Scheduler heartbeat	✗
3	DB connection reset	✓
4	Memory warning	✓
5	Debug trace	✗

There are 4 true relevant logs in total.

Precision@K — “Purity”

How many of the retrieved docs are relevant?

$$Precision@5 = \frac{\text{relevant retrieved}}{\text{total retrieved}} = 3/5 = 0.6$$

✓ 60% of what we retrieved is useful.

⚠ Low Precision = noisy retrieval → wasted tokens.

Recall@K — “Coverage”

How many of the relevant docs did we actually find?

$$Recall@5 = \frac{\text{relevant retrieved}}{\text{total relevant}} = 3/4 = 0.75$$

✓ We covered 75% of the right info.

⚠ Low Recall = missing facts → model will hallucinate.

MRR (Mean Reciprocal Rank)

How early do the correct results appear?

If the first relevant doc appears at rank 2 $\rightarrow$ Reciprocal Rank = $1/2 = 0.5$

Average across queries to get **MRR**.

✅ Higher MRR $\rightarrow$ users find answers faster.

nDCG (Normalized Discounted Cumulative Gain)

The most powerful retrieval metric.

It rewards ranking quality — relevant documents higher in the list matter more.

$$DCG@K = \sum_{i=1}^K \frac{2^{rel_i} - 1}{\log_2(i + 1)}$$
$$nDCG@K = \frac{DCG@K}{IDCG@K}$$

Example:

Rank	rel	Contribution
1	3	$7/1 = 7.00$
2	2	$3/1.58 = 1.89$
3	1	$1/2 = 0.5$

$DCG@5 = 9.39$

If perfect $\rightarrow IDCG@5 = 9.39$

$nDCG = 1.0$ ✅ perfect ranking.

If order shuffled $\rightarrow DCG = 6.9 \rightarrow nDCG = 0.74$

✅ nDCG shows how close your ranking is to ideal.

Understanding nDCG — Normalized Discounted Cumulative Gain

nDCG measures how good your ranking is compared to the *best possible (ideal) ranking*.

It's the gold-standard metric for **search engines, recommendation systems, and RAG retrievers**.

⚙️ Step 1: Formula

Discounted Cumulative Gain (DCG):

$$\text{DCG}@K = \sum ((2^{\text{rel}_i} - 1) / \log_2(i + 1))$$

where

- rel_i = relevance score of item i
- i = rank position (1-based)
- K = number of results you evaluate

Normalized DCG (nDCG):

$$\text{nDCG}@K = \text{DCG}@K / \text{IDCG}@K$$

$\text{IDCG}@K$ is the ideal DCG if results were perfectly ranked.

⚙️ Step 2: Example — Perfect Ranking

rel (Relevance)

How relevant the item is (usually human-judged)

- 3 = highly relevant

- 2 = relevant
- 1 = somewhat relevant
- 0 = not relevant

Rank	rel	Contribution Formula	Contribution
1	3	$(2^3-1)/\log_2(1+1) = 7/1$	7.00
2	2	$(2^2-1)/\log_2(2+1) = 3/1.58$	1.89
3	1	$(2^1-1)/\log_2(3+1) = 1/2$	0.50

$$\text{DCG@3} = 7 + 1.89 + 0.50 = 9.39$$

$$\text{IDCG@3} = 9.39$$

$$\text{nDCG@3} = 9.39 / 9.39 = 1.0 \quad \checkmark \text{ perfect ranking}$$

⚙️ Step 3: Example — Shuffled Ranking

Suppose the ranking is wrong:

Rank	rel
1	1
2	3
3	2

Compute DCG:

$$(2^1-1)/\log_2(1+1) = 1/1 = 1.00$$

$$(2^3-1)/\log_2(2+1) = 7/1.58 = 4.43$$

$$(2^2-1)/\log_2(3+1) = 3/2 = 1.50$$

$$\text{DCG@3} = 1 + 4.43 + 1.50 = 6.93$$

$$\text{IDCG@3} = 9.39$$

$$\text{nDCG@3} = 6.93 / 9.39 = 0.74$$

✅ **nDCG = 0.74 → 74% as good as the ideal ranking**

✅ **Key Takeaways**

- **nDCG = 1.0 → perfect ranking**
- **nDCG ≈ 0 → poor ranking**
- High-relevance items at top positions boost the score exponentially
- Log discount mimics human behavior — users care more about top results

🧩 **Retrieval Metrics Summary**

🧩 Retrieval Metrics Summary

Metric	Focus	Ideal
Precision@K	% of top K results relevant	≥ 0.7
Recall@K	% of all relevant results retrieved	≥ 0.8
MRR	Early relevance	≥ 0.5
nDCG	Overall ranking quality	≥ 0.85

🔥 **PART 2 — Generation Metrics**

How good is your “writer”?

The retriever feeds the context; now the LLM must produce an answer that’s grounded, relevant, and complete.

Generation Precision — “Groundedness”

How much of what the model said is actually supported by the context?

$$GenerationPrecision = \frac{\text{supported statements}}{\text{total statements}}$$

Example:

LLM generated 10 sentences → 8 supported → **Precision = 0.8**

✅ Higher = fewer hallucinations.

⚠️ Low = model is making things up.

Generation Recall — “Completeness”

Of all relevant facts available, how many did the LLM actually use?

$$GenerationRecall = \frac{\text{relevant context used}}{\text{total relevant context}}$$

Example:

10 relevant facts in retrieved chunks → LLM used 7 → **Recall = 0.7**

✅ Higher = more comprehensive.

⚠️ Low = answer might miss key details.

Faithfulness

Checks that all statements are consistent with context (no contradictions).

If 9 of 10 statements can be verified → Faithfulness = 0.9 ✅

Context Precision — “Did we give the model too much junk?”

How efficiently the LLM used the retrieved context.

$$\text{ContextPrecision} = \frac{\text{used chunks}}{\text{retrieved chunks}}$$

Retriever fetched 5 chunks; LLM used 3 → 0.6

⚠ Low = too many irrelevant chunks (optimize retriever).

Answer Relevance — “Did it answer the question?”

How semantically aligned is the generated answer with the user’s query?

$$\text{AnswerRelevance} = \text{cosine_sim}(\text{query}, \text{answer})$$

Use cosine similarity or LLM scoring:

Example:

LLM output = “It failed due to timeout.”

Ground Truth = “Timeout & DB reset.”

→ Relevance = 0.9 

If output = “Airflow is a scheduler.” → 0.1 

Generation Metric Summary

🧠 Generation Metric Summary

Metric	Focus	What It Measures	Ideal
Generation Precision	Faithfulness	% statements grounded in context	≥ 0.8
Generation Recall	Completeness	% of context facts used	≥ 0.7
Context Precision	Efficiency	% of retrieved chunks used	≥ 0.6
Answer Relevance	Intent match	Query ↔ Answer alignment	≥ 0.85

🤖 PART 3 — Agentic RAG Metrics

When your AI starts reasoning and acting

Agentic RAG introduces autonomous reasoning and tool execution, like CrewAI or LangGraph agents calling APIs or validating results.

So, you add a new metric layer: Action-level metrics.

🎯 Tool Accuracy

How many agentic actions (like database queries or log fetches) succeeded?

$$\text{Tool Accuracy} = \frac{\text{successful actions}}{\text{total actions}}$$

If 9 out of 10 tool calls worked → 0.9

🎯 Reasoning Trace Consistency

How consistent was the agent's step-by-step reasoning with retrieved evidence?

A banner for Medium dot calm. The background is a soft-focus image of a green field under a blue sky with light clouds. The text 'Medium dot calm.' is in a large, bold, black sans-serif font. Below it, 'Longer reads. Deeper breaths.' is in a smaller, regular black sans-serif font.

Medium dot calm.

Longer reads. Deeper breaths.

Upgrade today

LLM evaluators (or humans) check:

- Did the agent cite facts correctly?
- Were actions logically sequenced?

Typical score range: $0 \rightarrow 1$

Multi-Hop Answer Faithfulness

In Agentic RAG, the model might:

Retrieve data

Query another tool

Generate the final answer

Each hop compounds risk.

You measure **Faithfulness per hop**, then aggregate.

Example:

Hop 1 (Retrieval): 0.9

Hop 2 (Computation): 0.8

Hop 3 (Final Answer): 0.85

→ Avg Faithfulness = 0.85 

PART 4 — RAGAS (Comprehensive Evaluation)

RAGAS (Retrieval-Augmented Generation Assessment Suite)
combines all the above into one interpretable score.

$$RAGAS = 0.4 \times Faithfulness + 0.2 \times ContextPrecision + 0.2 \times ContextRecall + 0.2 \times AnswerRelevance$$

Example:

Faithfulness = 0.9

ContextPrecision = 0.7

ContextRecall = 0.8

AnswerRelevance = 0.9

→ RAGAS = 0.84  Production-ready

 How These Metrics Differ Between RAG and Agentic RAG

Layer	RAG	Agentic RAG
Retrieval Metrics	Precision, Recall, nDCG	Same
Generation Metrics	Faithfulness, Answer Relevance	Same + Reasoning validation
Action Metrics	N/A	Tool accuracy, multi-hop consistency
System Metrics	RAGAS	Extended RAGAS (includes agent hops)
Optimization Goal	Accuracy & grounding	Accuracy + autonomy + safety

 RAG = Knows the right thing

 Agentic RAG = Knows + Does the right thing

 PART 5 — Production Integration

You can integrate all these metrics in your MLOps stack using:

Tool	Purpose
RAGAS	Grounding & relevance
TruLens	Real-time evaluation dashboard
LangSmith	Traces & dataset analytics
Evidently AI + Prometheus	Drift & stability monitoring
Grafana	System metrics visualization

Example: Evaluating an Incident Copilot

Metric	Value	Target	Meaning
Recall@5	0.84	≥ 0.8	Key logs retrieved
Precision@5	0.68	≥ 0.7	Some noise
Faithfulness	0.9	≥ 0.8	Grounded output
Generation Recall	0.75	≥ 0.7	Sufficient completeness
Answer Relevance	0.91	≥ 0.85	Aligned with query
RAGAS	0.84	≥ 0.8	Overall strong
MTTR Reduction	60%	—	Proven business impact

Key Takeaways

- **Retrieval metrics** (Precision, Recall, nDCG) tell you *if the right facts were found*.
- **Generation metrics** (Faithfulness, Context Precision, Answer Relevance) tell you *if the LLM used and explained them correctly*.
- **Agentic RAG metrics** add *tool accuracy and reasoning consistency*.
- **RAGAS** ties it all together into a single, trackable score.

- ✓ High precision → trustworthy answers
- ✓ High recall → complete answers
- ✓ High relevance → helpful answers
- ✓ High faithfulness → safe answers

A great GenAI system doesn't just answer questions — it proves it knows *why* its answer is correct.

✓ Official RAGAS Metrics

1. context_precision

Definition (RAGAS):

Measures the proportion of retrieved contexts that are relevant to the ground truth answer.

- Evaluates retriever quality.
- Higher = fewer irrelevant chunks retrieved.
- Does NOT measure how the LLM “used” the context.

✓ Retriever metric

✗ Not a generation metric

2. context_recall

Definition (RAGAS):

Measures whether all relevant information needed to answer the question was retrieved.

- Higher = retrieval covers necessary facts.
- Lower = retriever missed important context.

✓ Retriever metric

3. faithfulness

Definition (RAGAS):

Measures whether the generated answer is fully supported by the retrieved context.

- Detects hallucinations.
- Checks grounding of answer to context.

✓ Generation metric

✓ Groundedness metric

4. answer_correctness

Definition (RAGAS):

Measures semantic similarity between the generated answer and the ground truth answer.

- Evaluates factual correctness.
- Independent of whether hallucination occurred.

✓ Generation metric

5. answer_relevancy **(sometimes called answer relevance)**

Definition (RAGAS):

Measures whether the generated answer addresses the user's question.

- Detects off-topic responses.

✓ Generation metric

What Is NOT Official in RAGAS

These are NOT official RAGAS metrics:

- “Generation Precision”
- “Generation Recall”
- “LLM used 3 out of 5 chunks”

Those are conceptual explanations, not implemented RAGAS metrics.

Correct Categorization

Retriever Metrics

- `context_precision`
- `context_recall`

Generation Metrics

- `faithfulness`
- `answer_correctness`
- `answer_relevancy`

Perfect — let’s compare them cleanly and formally.

Precision@K vs RAGAS `context_precision`

Aspect	Precision@K (IR)	RAGAS <code>context_precision</code>
Field	Information Retrieval	RAG Evaluation
Evaluates	Retriever	Retriever (semantic)
Requires manual relevance labels?	✅ Yes	❌ No
Uses ground-truth answer?	❌ No	✅ Yes
Uses LLM / embeddings?	❌ No	✅ Yes
Mathematical formula	Relevant@K / K	Semantically relevant contexts / K
Measures	Retrieval purity	Retrieval usefulness for answering
Deterministic?	Yes (given labels)	No (depends on semantic judge)

1 Precision@K (Classic IR)

Definition

$$Precision@K = \frac{\text{Number of relevant documents in top } K}{K}$$

You must already know which documents are relevant.

Example

Top 5 retrieved docs:

- 3 are labeled relevant
- 2 are irrelevant

[Precision@5 = 3/5 = 0.6]

- ✓ Pure retriever metric
- ✓ No generation involved

Definition (Conceptual)

Measures how many retrieved contexts are relevant to producing the **ground-truth answer**.

Instead of manual labels, RAGAS:

- Uses semantic similarity
- Uses embeddings and/or LLM judgment
- Evaluates relevance relative to the ground truth answer

Example

Top 5 retrieved chunks:

- 3 contain information needed for ground truth
- 2 are unrelated

Score ≈ 0.6

But the relevance is determined semantically, not manually.

Key Conceptual Difference

Precision@K asks:

“Is this document relevant to the question?”

RAGAS context_precision asks:

“Was this retrieved context useful for generating the correct answer?”

Subtle but important difference.

When They Match

If:

- You have perfect manual labels
- RAGAS semantic judge is accurate

Then:

```
Precision@K  $\approx$  context_precision
```

But they are computed differently.

When They Differ

They may differ when:

- A document is related to the question but not needed for the answer.
- A document partially overlaps semantically.
- Semantic similarity model misjudges relevance.

Clean Mental Model

Precision@K = Hard-labeled document relevance

RAGAS context_precision = Soft semantic answer-support relevance



Follow

Written by Dewasheesh Rana

505 followers · 2 following

Dewasheesh Rana is an AI/ML Architect with 14+ yrs experience in Agentic AI, LLMs, RAG, MLOps, API's building scalable systems on Azure, AWS and PCF at scale.

No responses yet



Ramesh Chukka

What are your thoughts?



